

▼ Telco Customer Churn – EDA and Business Understanding

This notebook performs exploratory data analysis (EDA) for a telecom customer churn dataset. The goal is to understand the business problem, explore the data, and identify key patterns that influence whether a customer leaves the company (churns) or stays.

Business Problem

In the telecom industry, customers can switch providers easily, and losing existing customers (customer churn) directly reduces recurring revenue. Acquiring a new customer is typically more expensive than retaining an existing one, so telecom operators focus heavily on predicting and reducing churn. [web:47][web:49]

The objective of this project is to build a data-driven churn prediction system that:

- Identifies customers who are likely to churn in the near future.
- Highlights key drivers of churn so the business can design targeted retention campaigns.
- Provides interpretable insights that can guide pricing, service quality, and contract strategies.

▼ 1. Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

▼ 2. Load raw data

```
df = pd.read_csv("Telco-Customer-Churn-dataset.csv")
df.head()
df.info()
df.describe(include="all")
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   customerID  7043 non-null   object  
 1   gender       7043 non-null   object  
 2   SeniorCitizen 7043 non-null   int64  
 3   Partner      7043 non-null   object  
 4   Dependents   7043 non-null   object  
 5   tenure       7043 non-null   int64  
 6   PhoneService 7043 non-null   object  
 7   MultipleLines 7043 non-null   object  
 8   InternetService 7043 non-null   object  
 9   OnlineSecurity 7043 non-null   object  
 10  OnlineBackup  7043 non-null   object  
 11  DeviceProtection 7043 non-null   object  
 12  TechSupport   7043 non-null   object  
 13  StreamingTV  7043 non-null   object  
 14  StreamingMovies 7043 non-null   object  
 15  Contract     7043 non-null   object  
 16  PaperlessBilling 7043 non-null   object  
 17  PaymentMethod 7043 non-null   object  
 18  MonthlyCharges 7043 non-null   float64 
 19  TotalCharges  7043 non-null   object  
 20  Churn        7043 non-null   object  
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

2. Dataset Description

The dataset contains customer-level information from a fictional telecom company that provides phone and internet services. Each row represents one customer and includes: [web:35][web:57]

- **Customer demographics** - e.g. gender, senior citizen flag, partner, dependents.
- **Account Information** - e.g. tenure (months with the company), contract type, payment method, paperless billing.
- **Services subscribed** - e.g. phone service, multiple lines, internet service type, online security, tech support, streaming services.
- **Billing Information** - e.g. monthly charges, total charges.
- **Target variable: Churn** indicates whether the customer left the company in the last month (**Yes**) or not (**No**). [web:35][web:57]

In this notebook, the focus is on understanding how these features relate to churn, not on building models yet.

3. Basic Cleaning (TotalCharges)

dtypes: float64(1), int64(2), object(18)

memory usage: 1.1+ MB

Data Cleaning (TotalCharges Fix)

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...
count	7043	7043	7043.000000	7043	7043	7043.000000	7043	7043	7043	7043	7043

During inspection, the **TotalCharges** column appears as an object (string) type due to spaces in some rows. This prevents correct numeric analysis. To address this, the notebook:

- **top** 3186-AJIEK Male NaN No No NaN Yes No Fiber optic No
- Replaces blank spaces in **TotalCharges** with **NaN**.
- **freq** Converts **TotalCharges** to a numeric type. 3641 4933 NaN 6361 3390 3096 3498
- **dropna** Drops rows where **TotalCharges** is missing after conversion. NaN 32.371149 NaN NaN NaN NaN NaN

This ensures that billing information can be used correctly in later analysis and modeling.

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...
min	NaN	NaN	0.000000	NaN	NaN	0.000000	NaN	NaN	NaN	NaN	NaN

```
df["TotalCharges"] = df["TotalCharges"].replace(" ", np.nan)
df["TotalCharges"] = pd.to_numeric(df["TotalCharges"])
df = df.dropna(subset=["TotalCharges"])
df = df.reset_index(drop=True)

df.info()
df.describe()
11 rows x 21 columns
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7032 entries, 0 to 7031
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customerID      7032 non-null    object  
 1   gender          7032 non-null    object  
 2   SeniorCitizen   7032 non-null    int64  
 3   Partner         7032 non-null    object  
 4   Dependents     7032 non-null    object  
 5   tenure          7032 non-null    int64  
 6   PhoneService    7032 non-null    object  
 7   MultipleLines   7032 non-null    object  
 8   InternetService 7032 non-null    object  
 9   OnlineSecurity  7032 non-null    object  
 10  OnlineBackup    7032 non-null    object  
 11  DeviceProtection 7032 non-null    object  
 12  TechSupport    7032 non-null    object  
 13  StreamingTV    7032 non-null    object  
 14  StreamingMovies 7032 non-null    object  
 15  Contract        7032 non-null    object  
 16  PaperlessBilling 7032 non-null    object  
 17  PaymentMethod   7032 non-null    object  
 18  MonthlyCharges 7032 non-null    float64 
 19  TotalCharges   7032 non-null    float64 
 20  Churn          7032 non-null    object  
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB

```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	
count	7032.000000	7032.000000	7032.000000	7032.000000	
mean	0.162400	32.421786	64.798208	2283.300441	
std	0.368844	24.545260	30.085974	2266.771362	
min	0.000000	1.000000	18.250000	18.800000	
25%	0.000000	9.000000	35.587500	401.450000	
50%	0.000000	29.000000	70.350000	1397.475000	
75%	0.000000	55.000000	89.862500	3794.737500	
max	1.000000	72.000000	118.750000	8684.800000	

4. Target distribution

Target Variable: Churn Distribution

Understanding the distribution of the target variable (`Churn`) is essential: [web:25][web:52]

- A **countplot** shows how many customers churned vs. did not churn.
- A **percentage view** (pie chart) highlights class imbalance.

In many telecom datasets, churers are a minority compared to non-churers. This imbalance affects model choice and evaluation, because accuracy alone can be misleading if most customers do not churn.

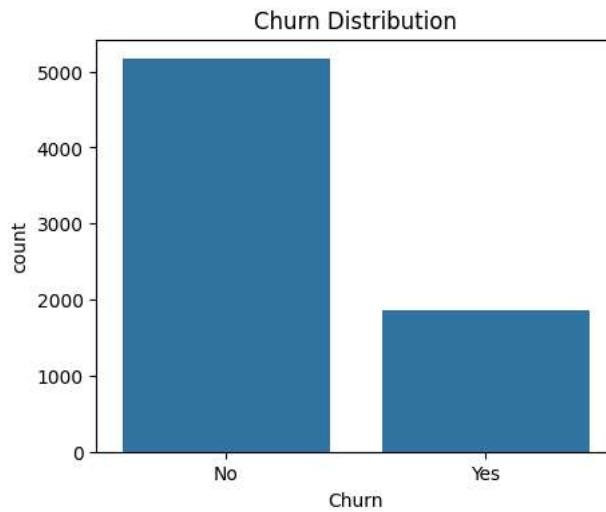
```

plt.figure(figsize=(5, 4))
sns.countplot(x="Churn", data=df)
plt.title("Churn Distribution")
plt.show()

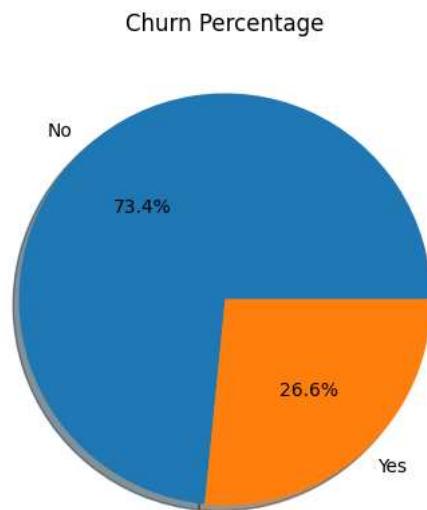
print(df["Churn"].value_counts(normalize=True) * 100)

plt.figure(figsize=(5, 5))
df["Churn"].value_counts().plot.pie(
    autopct="%1.1f%%", labels=["No", "Yes"], shadow=True
)
plt.title("Churn Percentage")
plt.ylabel("")
plt.show()

```



```
Churn
No      73.421502
Yes     26.578498
Name: proportion, dtype: float64
```



5. Numerical features – distributions

Numerical Features – Univariate Analysis

Key numeric variables in this dataset include:

- `tenure` – number of months the customer has stayed with the company.
- `MonthlyCharges` – amount charged to the customer each month.
- `TotalCharges` – total amount billed over the customer's lifetime.

Using histograms and KDE plots, the notebook explores:

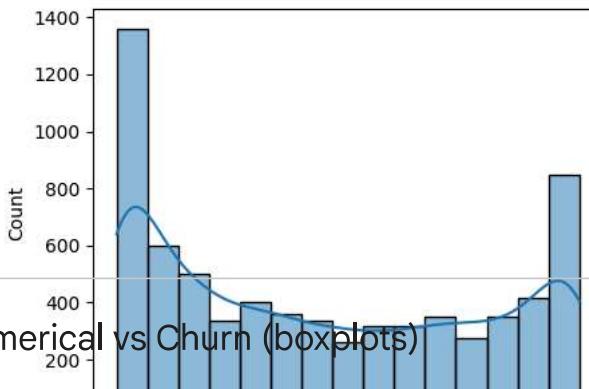
- Overall distribution and spread of each variable.
- Presence of skewness or outliers.

These insights help later when choosing transformations, scaling, and model types.

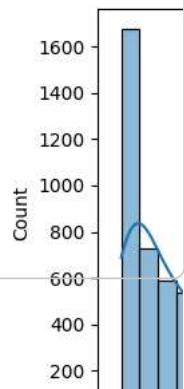
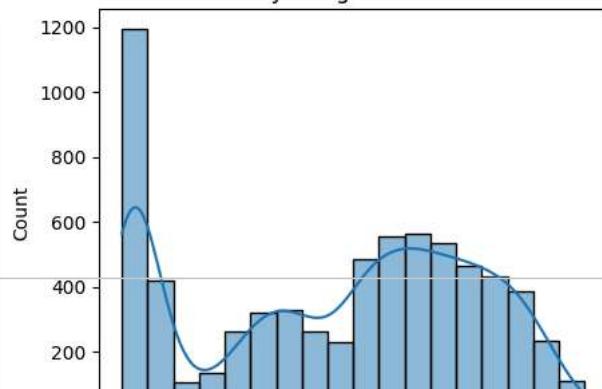
```
num_cols = ["tenure", "MonthlyCharges", "TotalCharges"]

plt.figure(figsize=(14, 4))
for i, col in enumerate(num_cols, 1):
    plt.subplot(1, 3, i)
    sns.histplot(df[col], kde=True)
    plt.title(f"{col} Distribution")
plt.tight_layout()
plt.show()
```

tenure Distribution

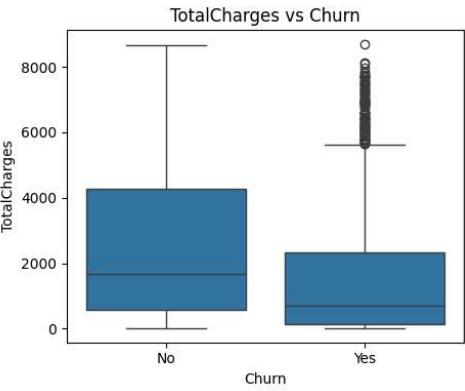
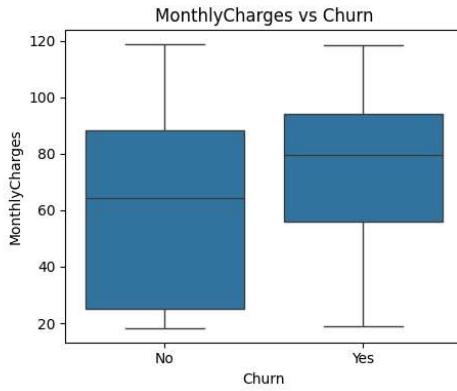
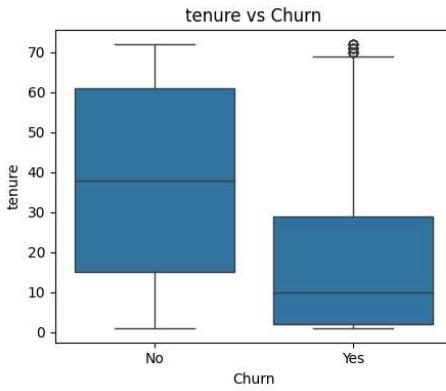


MonthlyCharges Distribution



Numerical vs Churn (boxplots)

```
plt.figure(figsize=(14, 4))
for i, col in enumerate(num_cols, 1):
    plt.subplot(1, 3, i)
    sns.boxplot(x="Churn", y=col, data=df)
    plt.title(f"{col} vs Churn")
plt.tight_layout()
plt.show()
```



Numerical Features vs Churn

To see how numeric features relate to churn, the notebook uses boxplots and density plots split by churn label.

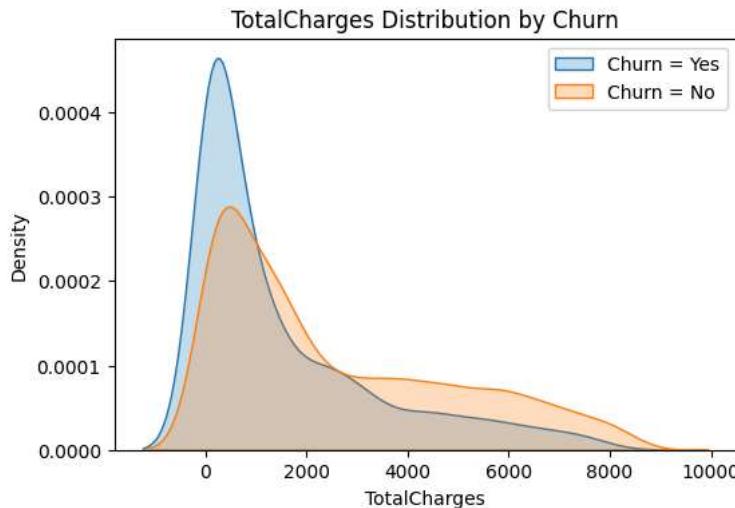
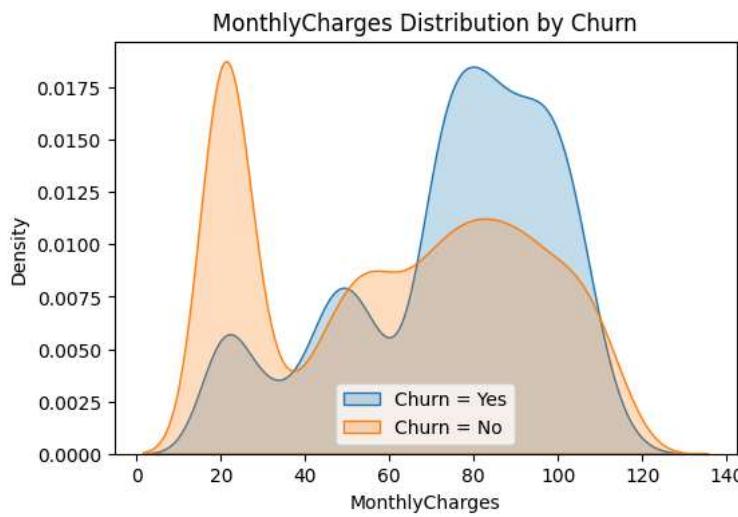
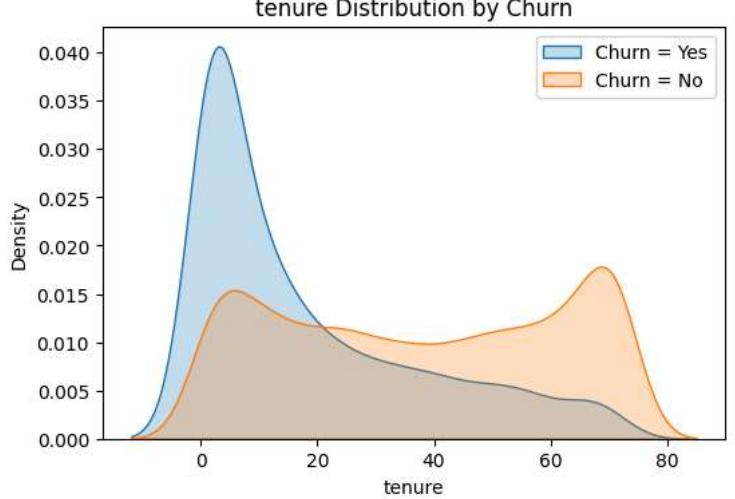
Typical patterns in telecom churn datasets include:

- **Tenure:** customers with shorter tenure tend to churn more, while long-tenure customers are more loyal.
- **MonthlyCharges:** higher monthly charges can be associated with higher churn, as customers are more price-sensitive.
- **TotalCharges:** chunbers often have lower total charges because they leave earlier in their lifecycle. [web:47][web:49]

These relationships provide early evidence for which variables are strong churn drivers.

KDE per churn class

```
for col in num_cols:
    plt.figure(figsize=(6, 4))
    sns.kdeplot(
        df[df["Churn"] == "Yes"][col],
        label="Churn = Yes",
        fill=True,
        common_norm=False,
    )
    sns.kdeplot(
        df[df["Churn"] == "No"][col],
        label="Churn = No",
        fill=True,
        common_norm=False,
    )
    plt.title(f"{col} Distribution by Churn")
    plt.legend()
    plt.show()
```



6. Categorical vs Churn

Categorical Features vs Churn

This section examines how different categorical variables behave for churners vs non-churners using grouped bar charts. Important groups include: [web:47][web:52]

- **Contract type** (Month-to-month, One-year, Two-year).
- **Payment method** (Electronic check vs other methods).
- **Internet service type** (DSL, Fiber optic, No internet).
- **Security and support services** (OnlineSecurity, TechSupport, etc.).
- **Billing type** (PaperlessBilling).

Common patterns observed in telecom churn studies:

- Month-to-month contracts have significantly higher churn than long-term contracts.
- Customers paying via electronic check churn more than those using automatic or bank transfers.
- Fiber optic customers often show higher churn due to higher pricing or performance expectations.

- Customers without online security or tech support tend to churn more frequently.

```

cat_cols = [
    "Contract",
    "PaymentMethod",
    "InternetService",
    "OnlineSecurity",
    "TechSupport",
    "PaperlessBilling",
]

plt.figure(figsize=(18, 12))
for i, col in enumerate(cat_cols, 1):
    plt.subplot(3, 2, i)
    sns.countplot(x=col, hue="Churn", data=df)
    plt.title(f"{col} vs Churn")
    plt.xticks(rotation=30, ha="right")
plt.tight_layout()
plt.show()

```

