

## 02\_Preprocessing\_and\_Modeling

### ▼ Telco Customer Churn – Preprocessing & Model Comparison

This notebook builds a **production-ready sklearn pipeline** and compares multiple ML algorithms for churn prediction. Key objectives:  
[web:36][web:79]

1. **Design reproducible preprocessing** (ColumnTransformer for numeric + categorical features)
2. **Train & evaluate 5+ models** (Logistic Regression, Random Forest, XGBoost, CatBoost)
3. **Hyperparameter tuning** on top performers
4. **Select best model** based on business metrics (ROC-AUC, Recall for churners)
5. **Save complete pipeline** (`best_pipeline.pkl`) for deployment

### ▼ 1: Imports

```
import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    roc_auc_score,
    confusion_matrix,
    classification_report,
)
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier

!pip install catboost

from xgboost import XGBClassifier
from catboost import CatBoostClassifier
from sklearn.model_selection import RandomizedSearchCV

import matplotlib.pyplot as plt
import seaborn as sns
sns.set(style="whitegrid")

Requirement already satisfied: catboost in /usr/local/lib/python3.12/dist-packages (1.2.8)
Requirement already satisfied: graphviz in /usr/local/lib/python3.12/dist-packages (from catboost) (0.21)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (from catboost) (3.10.0)
Requirement already satisfied: numpy<3.0,>1.16.0 in /usr/local/lib/python3.12/dist-packages (from catboost) (2.0.2)
Requirement already satisfied: pandas>=0.24 in /usr/local/lib/python3.12/dist-packages (from catboost) (2.2.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (from catboost) (1.16.3)
Requirement already satisfied: plotly in /usr/local/lib/python3.12/dist-packages (from catboost) (5.24.1)
Requirement already satisfied: six in /usr/local/lib/python3.12/dist-packages (from catboost) (1.17.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (1.4.9)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (25.0)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (3.2.5)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.12/dist-packages (from plotly->catboost) (9.1.2)
```

### ▼ 2: Load cleaned data

#### ▼ Load Cleaned Data

Loading the cleaned dataset from `01_EDA_and_Business_Understanding.ipynb` [web:36]

- **Source:** `telco_clean.csv` (TotalCharges fixed, no missing values)
- **Target:** `Churn` mapped to binary (0=No, 1=Yes)
- **Features:** 20 columns (3 numeric + 17 categorical)
- **Split:** 80/20 train-test with stratification to preserve churn ratio

This ensures the modeling phase starts with clean, analysis-ready data.

```
df = pd.read_csv("telco_clean.csv")

# Drop ID column if present
if "customerID" in df.columns:
    df = df.drop(columns=["customerID"])

# Target
y = df["Churn"].map({"No": 0, "Yes": 1})
X = df.drop(columns=["Churn"])
```

## ✗ 3: Train-test split

### ✗ Train-Test Split

**Stratified split** (80/20) preserves the churn class distribution:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

print(X_train.shape, X_test.shape)

(5625, 19) (1407, 19)
```

## ✗ 4: Identify numeric & categorical features

### ✗ Feature Engineering Strategy

**Automatic feature type detection** for robust preprocessing pipeline: [web:36][web:79]

**Numeric features** (3) – will be **StandardScaled**:

**Categorical features** (17) – will be **OneHotEncoded**:

**ColumnTransformer** handles both automatically, producing ~45 features after one-hot encoding.

```
numeric_features = ["tenure", "MonthlyCharges", "TotalCharges"]
categorical_features = [c for c in X.columns if c not in numeric_features]
print("Numeric:", numeric_features)
print("Categorical:", categorical_features)

Numeric: ['tenure', 'MonthlyCharges', 'TotalCharges']
Categorical: ['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod']
```

## ✗ 5: Preprocessor

### ✗ Production-Ready Preprocessing Pipeline

**ColumnTransformer** design for end-to-end reproducibility:

**Key design choices:**

- `handle_unknown="ignore"` → Production-safe (new categories don't crash)
- `StandardScaler` → Tree models don't need it, but Logistic Regression does
- **Single pipeline** → Zero data leakage, same transform for train + test + production

```
numeric_transformer = Pipeline(steps=[("scaler", StandardScaler())])

categorical_transformer = Pipeline(steps=[("onehot", OneHotEncoder(handle_unknown="ignore"))])

preprocessor = ColumnTransformer(transformers=[("num", numeric_transformer, numeric_features), ("cat", categorical_transformer, categorical_features)])
```

## ✗ 6: Helper – evaluation function

### ✗ Evaluation Framework

**Custom evaluation function** computes business-relevant metrics:

Metric	Why it matters for churn
Accuracy	Overall correctness
Precision	Of predicted churners, what % actually churn? (avoid wasting retention budget)

**Metric**      **Why it matters for churn**

- Recall** Of actual churners, what % do we catch? (business priority #1)  
**F1-Score** Balance of precision + recall  
**ROC-AUC** Model's ability to rank customers by churn risk

**Business priority:** High **Recall** (catch most churners) > Precision (some false alarms OK).

```
def evaluate_model(name, model, X_test, y_test, results_list):  
    y_pred = model.predict(X_test)  
    y_prob = model.predict_proba(X_test)[:, 1]  
  
    acc = accuracy_score(y_test, y_pred)  
    prec = precision_score(y_test, y_pred)  
    rec = recall_score(y_test, y_pred)  
    f1 = f1_score(y_test, y_pred)  
    roc = roc_auc_score(y_test, y_prob)  
  
    print(f"== {name} ==")  
    print("Accuracy:", acc)  
    print("Precision:", prec)  
    print("Recall:", rec)  
    print("F1-Score:", f1)  
    print("ROC-AUC:", roc)  
    print("Confusion matrix:\n", confusion_matrix(y_test, y_pred))  
    print("\nClassification report:\n", classification_report(y_test, y_pred))  
  
    results_list.append(  
        {  
            "Model": name,  
            "Accuracy": acc,  
            "Precision": prec,  
            "Recall": rec,  
            "F1": f1,  
            "ROC_AUC": roc,  
        }  
    )
```

## 7: Baseline Logistic Regression

```
baseline_pipe = Pipeline(steps=[("preprocessor", preprocessor), ("model", LogisticRegression(max_iter=500, solver="lbfgs", class_weight="balanced"))])  
baseline_pipe.fit(X_train, y_train)  
  
results = []  
evaluate_model("Logistic Regression (baseline)", baseline_pipe, X_test, y_test, results)  
  
==== Logistic Regression (baseline) ====  
Accuracy: 0.7256574271499645  
Precision: 0.4901315789473684  
Recall: 0.7967914438502673  
F1-Score: 0.6069246435845214  
ROC-AUC: 0.8350930005021457  
Confusion matrix:  
[[723 310]  
 [ 76 298]]  
  
Classification report:  
precision recall f1-score support  
0 0.90 0.70 0.79 1033  
1 0.49 0.80 0.61 374  
  
accuracy 0.73 1407  
macro avg 0.70 0.75 0.70 1407  
weighted avg 0.79 0.73 0.74 1407
```

## 8: Random Forest

```
rf_pipe = Pipeline(steps=[("preprocessor", preprocessor), ("model", RandomForestClassifier(n_estimators=300, random_state=42, class_weight="balanced"))])  
rf_pipe.fit(X_train, y_train)  
evaluate_model("Random Forest", rf_pipe, X_test, y_test, results)  
  
==== Random Forest ====  
Accuracy: 0.7874911158493249  
Precision: 0.631578947368421  
Recall: 0.48128342245989303  
F1-Score: 0.5462822458270106  
ROC-AUC: 0.8152957742104145  
Confusion matrix:  
[[928 105]  
 [194 180]]
```

```

Classification report:
      precision    recall  f1-score   support

       0          0.83     0.90      0.86     1033
       1          0.63     0.48      0.55      374

    accuracy                           0.79      1407
   macro avg       0.73     0.69      0.70      1407
weighted avg       0.78     0.79      0.78      1407

```

## ▼ 9: XGBoost

```

xgb_pipe = Pipeline(
    steps=[
        ("preprocessor", preprocessor),
        (
            "model",
            XGBClassifier(
                n_estimators=400,
                learning_rate=0.05,
                max_depth=5,
                subsample=0.8,
                colsample_bytree=0.8,
                eval_metric="logloss",
                n_jobs=-1,
                random_state=42,
            ),
        ),
    ],
)

xgb_pipe.fit(X_train, y_train)
evaluate_model("XGBoost", xgb_pipe, X_test, y_test, results)

==== XGBoost ====
Accuracy: 0.783226723525231
Precision: 0.6055045871559633
Recall: 0.5294117647058824
F1-Score: 0.5649072753209701
ROC-AUC: 0.8206472503636673
Confusion matrix:
[[904 129]
 [176 198]]

Classification report:
      precision    recall  f1-score   support

       0          0.84     0.88      0.86     1033
       1          0.61     0.53      0.56      374

    accuracy                           0.78      1407
   macro avg       0.72     0.70      0.71      1407
weighted avg       0.78     0.78      0.78      1407

```

## ▼ 10: CatBoost

```

cat_pipe = Pipeline(
    steps=[
        ("preprocessor", preprocessor),
        (
            "model",
            CatBoostClassifier(
                iterations=500,
                learning_rate=0.05,
                depth=6,
                loss_function="Logloss",
                eval_metric="AUC",
                verbose=0,
                random_seed=42,
            ),
        ),
    ],
)

cat_pipe.fit(X_train, y_train)
evaluate_model("CatBoost", cat_pipe, X_test, y_test, results)

==== CatBoost ====
Accuracy: 0.7796730632551528
Precision: 0.6012658227848101
Recall: 0.5080213903743316
F1-Score: 0.5507246376811594
ROC-AUC: 0.8247977180839774
Confusion matrix:
[[907 126]
 [184 190]]

```

Classification report:				
	precision	recall	f1-score	support
0	0.83	0.88	0.85	1033
1	0.60	0.51	0.55	374
accuracy			0.78	1407
macro avg	0.72	0.69	0.70	1407
weighted avg	0.77	0.78	0.77	1407

## ▼ 11: Collect and visualize results

```
import pandas as pd

results_df = pd.DataFrame(results)
results_df
```

	Model	Accuracy	Precision	Recall	F1	ROC_AUC	
0	Logistic Regression (baseline)	0.725657	0.490132	0.796791	0.606925	0.835093	
1	Random Forest	0.787491	0.631579	0.481283	0.546282	0.815296	
2	XGBoost	0.783227	0.605505	0.529412	0.564907	0.820647	
3	CatBoost	0.779673	0.601266	0.508021	0.550725	0.824798	

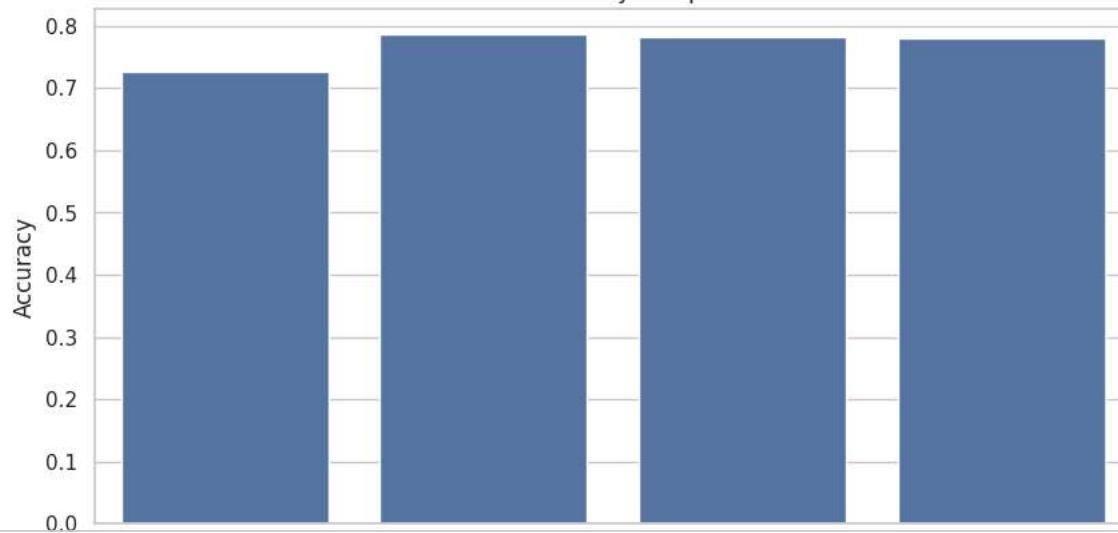
Next steps: [Generate code with results\\_df](#) [New interactive sheet](#)

## ▼ 12: Barplot of metrics

```
plt.figure(figsize=(10, 5))
sns.barplot(x="Model", y="Accuracy", data=results_df)
plt.xticks(rotation=30, ha="right")
plt.title("Model Accuracy Comparison")
plt.show()

plt.figure(figsize=(10, 5))
sns.barplot(x="Model", y="ROC_AUC", data=results_df)
plt.xticks(rotation=30, ha="right")
plt.title("Model ROC-AUC Comparison")
plt.show()
```

Model Accuracy Comparison



▼ 13: Simple hyperparameter tuning on best model (example: Random Forest)

```

param_dist = {
    "model__n_estimators": [200, 300, 500],
    "model__max_depth": [5, 10, 20, None],
    "model__min_samples_split": [2, 5, 10],
    "model__min_samples_leaf": [1, 2, 4],
}

rf_tuned = Pipeline(
    steps=[
        ("preprocessor", preprocessor),
        (
            "model",
            RandomForestClassifier(
                random_state=42,
                class_weight="balanced_subsample",
                n_jobs=-1,
            ),
        ),
    ],
)

rf_search = RandomizedSearchCV(
    rf_tuned,
    param_distributions=param_dist,
    n_iter=10,
    cv=3,
    scoring="roc_auc",
    n_jobs=-1,
    verbose=2,
    random_state=42,
)

rf_search.fit(X_train, y_train)
print("Best params:", rf_search.best_params_)
best_rf_pipe = rf_search.best_estimator_

evaluate_model("Random Forest (Tuned)", best_rf_pipe, X_test, y_test, results)
results_df = pd.DataFrame(results)
results_df

```

```

Fitting 3 folds for each of 10 candidates, totalling 30 fits
Best params: {'model__n_estimators': 200, 'model__min_samples_split': 2, 'model__min_samples_leaf': 4, 'model__max_depth': 10}
== Random Forest (Tuned) ==
Accuracy: 0.7555081734186212
Precision: 0.5277777777777778
Recall: 0.7620320855614974
F1-Score: 0.6236323851203501
ROC-AUC: 0.8329265262384105
Confusion matrix:
[[778 255]
 [ 89 285]]

Classification report:
precision    recall    f1-score   support
          0       0.90      0.75      0.82     1033
          1       0.53      0.76      0.62      374

   accuracy                           0.76      1407
  macro avg       0.71      0.76      0.72     1407
weighted avg       0.80      0.76      0.77     1407

          Model  Accuracy  Precision  Recall    F1  ROC_AUC
0  Logistic Regression (baseline)  0.725657  0.490132  0.796791  0.606925  0.835093
1           Random Forest  0.787491  0.631579  0.481283  0.546282  0.815296
2            XGBoost  0.783227  0.605505  0.529412  0.564907  0.820647
3            CatBoost  0.779673  0.601266  0.508021  0.550725  0.824798
4  Random Forest (Tuned)  0.755508  0.527778  0.762032  0.623632  0.832927

```

Next steps: [Generate code with results\\_df](#) [New interactive sheet](#)

## 14: Choose best model and save pipeline

```

best_idx = results_df["Recall"].idxmax()
best_name = results_df.loc[best_idx, "Model"]
print("Best model:", best_name)

if best_name == "Random Forest (Tuned)":
    final_model = best_rf_pipe
elif best_name == "Random Forest":
    final_model = rf_pipe
elif best_name == "XGBoost":
    final_model = xgb_pipe
elif best_name == "CatBoost":
    final_model = cat_pipe
else:
    final_model = baseline_pipe

import pickle

with open("best_pipeline.pkl", "wb") as f:
    pickle.dump(final_model, f)

results_df.to_csv("model_results.csv", index=False)

```